

03_1__CombineCategories

January 28, 2021

```
[1]: # The MIT License (MIT)

# Copyright (c) 2020, NVIDIA CORPORATION.

# Permission is hereby granted, free of charge, to any person obtaining a copy
# of
# this software and associated documentation files (the "Software"), to deal in
# the Software without restriction, including without limitation the rights to
# use, copy, modify, merge, publish, distribute, sublicense, and/or sell copies
# of
# the Software, and to permit persons to whom the Software is furnished to do
# so,
# subject to the following conditions:

# The above copyright notice and this permission notice shall be included in all
# copies or substantial portions of the Software.

# THE SOFTWARE IS PROVIDED "AS IS", WITHOUT WARRANTY OF ANY KIND, EXPRESS OR
# IMPLIED, INCLUDING BUT NOT LIMITED TO THE WARRANTIES OF MERCHANTABILITY,
# FITNESS
# FOR A PARTICULAR PURPOSE AND NONINFRINGEMENT. IN NO EVENT SHALL THE AUTHORS OR
# COPYRIGHT HOLDERS BE LIABLE FOR ANY CLAIM, DAMAGES OR OTHER LIABILITY, WHETHER
# IN AN ACTION OF CONTRACT, TORT OR OTHERWISE, ARISING FROM, OUT OF OR IN
# CONNECTION WITH THE SOFTWARE OR THE USE OR OTHER DEALINGS IN THE SOFTWARE
```

1 Tutorial: Feature Engineering for Recommender Systems

2 3. Feature Engineering - Categorical

2.1 3.1. Combining Categories / Cross Columns

2.2 Theory

Combining Categories (CC) is a simple, powerful technique, but often undervalued. We will use this strategy in other feature engineering techniques, as well, and will introduce its value in a simple example. In some datasets, categories by itself provide no information to predict the target. But if we combine multiple categories, together, then we can indentify patterns. For example, we have the following categories:

Weekday

Hour of the day

Each of them independently has no significant pattern in the dataset. If we combine them with *Weekday_HourOfTheDay*, then we can observe some strong behavior for certain times on the weekend. Decision Trees determine the split in the dataset on single features. If each categorical feature by itself does not provide the information gain, then Decision Trees cannot find a good split. If we provide a combined categorical feature, the Decision Tree can easier split the dataset.

Combining categories, also called Cross Column or Cross Product, is used in the [Wide Deep Architecture](#) by Google and is implemented in [Tensorflow](#)

```
[2]: import IPython

import cudf

import pandas as pd
import numpy as np
```

```
[3]: f1 = [0]*45 + [1]*45 + [2]*10 + [0]*5 + [1]*5 + [2]*90 + [0]*5 + [1]*5 + [2]*90
      ↪+ [0]*45 + [1]*45 + [2]*10
f2 = [0]*45 + [0]*45 + [0]*10 + [1]*5 + [1]*5 + [1]*90 + [0]*5 + [0]*5 + [0]*90
      ↪+ [1]*45 + [1]*45 + [1]*10
t = [1]*45 + [1]*45 + [1]*10 + [1]*5 + [1]*5 + [1]*90 + [0]*5 + [0]*5 + [0]*90
    ↪+ [0]*45 + [0]*45 + [0]*10

data = cudf.DataFrame({
    'f1': f1,
    'f2': f2,
})

for i in range(3,5):
    data['f' + str(i)] = np.random.choice(list(range(3)), data.shape[0])

data['target'] = t
```

```
[4]: data.head()
```

```
[4]:   f1  f2  f3  f4  target
0    0   0   0   0      1
1    0   0   1   1      1
2    0   0   0   2      1
3    0   0   0   0      1
4    0   0   1   2      1
```

We take a look on the features *f1* and *f2*. Each of the feature provides no information gain as each category has a 0.5 probability for the target.

```
[5]: data.groupby('f1').target.agg(['mean', 'count'])
```

```
[5]:      mean  count
f1
0      0.5    100
1      0.5    100
2      0.5    200
```

```
[6]: data.groupby('f2').target.agg(['mean', 'count'])
```

```
[6]:      mean  count
f2
0      0.5    200
1      0.5    200
```

If we analyze the features *f1* and *f2* together, we can observe a significant pattern in the target variable.

```
[7]: data.groupby(['f1', 'f2']).target.agg(['mean', 'count'])
```

```
[7]:      mean  count
f1 f2
0  0      0.9     50
   1      0.1     50
1  0      0.9     50
   1      0.1     50
2  0      0.1    100
   1      0.9    100
```

Next, we train a simple Decision Tree to show how combining categories will support the decision boundaries.

```
[8]: df = data.to_pandas()
```

```
[9]: import pydotplus
import sklearn.tree as tree
from IPython.display import Image
```

```
[10]: def get_hotn_features(df):
    out = []
    for col in df.columns:
        if col != 'target':
            out.append(pd.get_dummies(df[col], prefix=col))
    return pd.concat(out, axis=1)

def viz_tree(df, lf):
    dt_feature_names = list(get_hotn_features(df).columns)
    dt_target_names = 'target'
```

```

tree.export_graphviz(lf, out_file='tree.dot',
                     feature_names=dt_feature_names,
    ↪ class_names=dt_target_names,
                     filled=True)
graph = pydotplus.graph_from_dot_file('tree.dot')
return(graph.create_png())

```

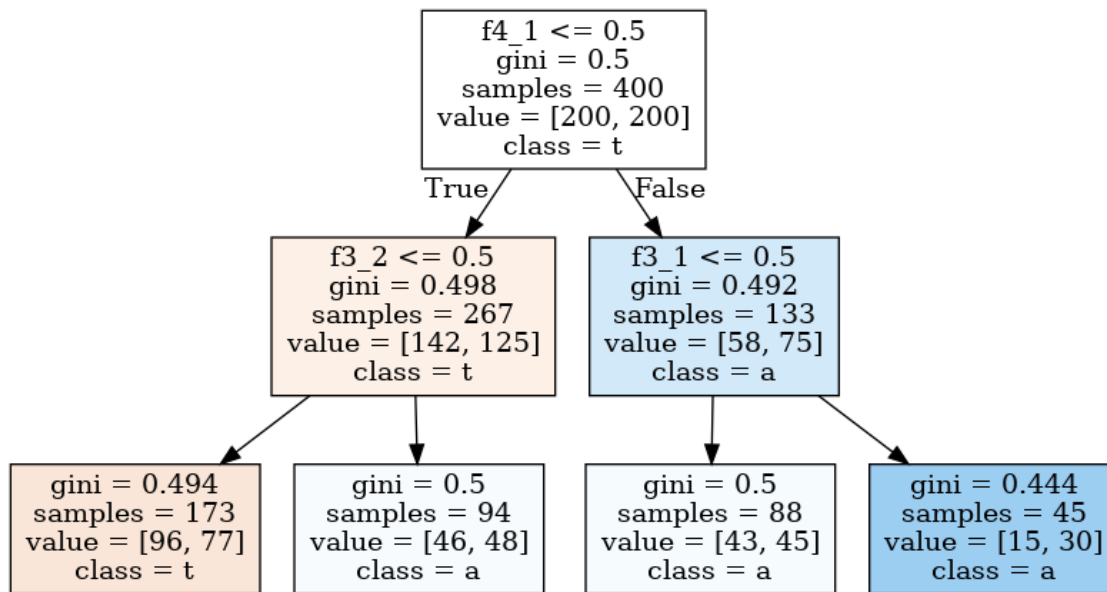
First, we train it without the combined categories $f1$ and $f2$. We can see, that the Decision Trees creates the split on the random features $f3$ and $f4$. The leaves have only a small information gain (e.g. 98 negative vs. 82 positive).

```

[11]: lf = tree.DecisionTreeClassifier(max_depth=2)
lf.fit(get_hotn_features(df), df[['target']])
Image(viz_tree(df, lf))

```

[11]:



Now, we combine the categories $f1$ and $f2$ as an additional feature. We can see that the Decision Tree uses that feature first and that the splits have a high information gain. For example, 190 negative vs. 110 positives.

```

[12]: df['f1_f2'] = df['f1'].astype(str) + df['f2'].astype(str)

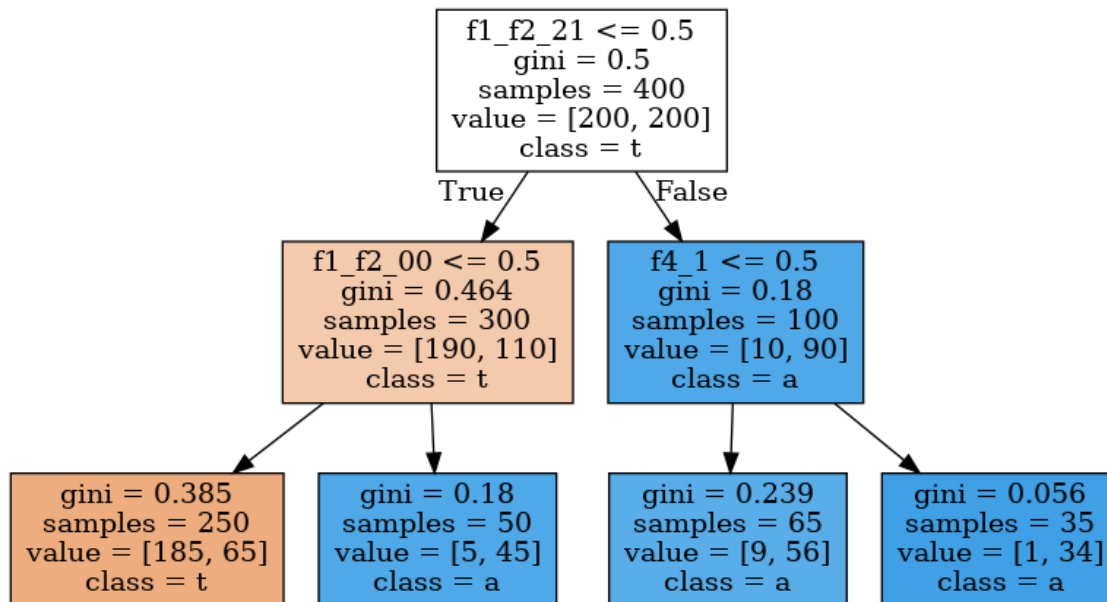
```

```

[13]: lf.fit(get_hotn_features(df), df[['target']])
Image(viz_tree(df, lf))

```

[13]:



This simple technique will be used in combination with other feature engineering techniques. We may have the idea - that is great, let's combine all categories into one feature. Unfortunately, this is not that easy. We want to balance the number of categories used, the number of observations in resulting category values and the information gain:

The more categories we combine, we will identify more underlying patterns - but combining more categories together reduces the number of observation per category in the resulting features

Higher number of observation in the resulting category shows a strong pattern and it is more generalizable

High information gain supports our model, but only if it is generalizable

The extreme example is that we combine all features f1, f2, f3 and f4 together. But the observation per category (count) is very small (4-20)

```
[14]: df.groupby([x for x in df.columns if 'target' not in x and 'f1_f2' not in x]).
      ↪target.agg(['mean', 'count']).head(10)
```

```
[14]:
```

					mean	count
	f1	f2	f3	f4		
	0	0	0	0	1.000000	6
				1	1.000000	5
				2	0.750000	4
		1	0		0.750000	4
			1		0.857143	7
			2		1.000000	5
		2	0		0.750000	4
			1		0.750000	4

		2	1.000000	11
1	0	0	0.000000	6

Best practice:

Combining low cardinal categories is a good start. For example, the dataset size is 100M rows and there are multiple categories with a cardinality (# of unique values) of 10-50, then combining them should not result in low observation count

Exploratory Data Analysis (EDA) is faster than training a model. Analyzing the information value for different combination of categorical features (on a sample) is really fast.

Example of getting the cardinality for categories:

```
[15]: df.astype(str).describe()
```

```
[15]:
```

	f1	f2	f3	f4	target	f1_f2
count	400	400	400	400	400	400
unique	3	2	3	3	2	6
top	2	0	1	0	0	21
freq	200	200	136	142	200	100

Summary Combining categories identifies underlying patterns in the dataset

The technique can support Decision Trees to create better splits as Decision Trees analyze features independently of each other

2.3 Practice

Now, it is your turn. What are good combinations of categories in our dataset?

ToDo:

Define which categorical features should be combined? Why should these be combined? What are your hypotheses?

```
[16]: import cudf
```

```
[17]: df_train = cudf.read_parquet('./data/train.parquet')
```

```
[18]: df_train.head()
```

```
[18]:
```

		event_time	event_type	product_id	brand	price	user_id	\
0	2019-12-01	00:00:28 UTC	cart	17800342	zeta	66.90	550465671	
1	2019-12-01	00:00:39 UTC	cart	3701309	polaris	89.32	543733099	
2	2019-12-01	00:00:40 UTC	cart	3701309	polaris	89.32	543733099	
3	2019-12-01	00:00:41 UTC	cart	3701309	polaris	89.32	543733099	
4	2019-12-01	00:01:56 UTC	cart	1004767	samsung	235.60	579970209	

	user_session	target	cat_0	cat_1	\
--	--------------	--------	-------	-------	---

0	22650a62-2d9c-4151-9f41-2674ec6d32d5	0	computers	desktop
1	a65116f4-ac53-4a41-ad68-6606788e674c	0	appliances	environment
2	a65116f4-ac53-4a41-ad68-6606788e674c	0	appliances	environment
3	a65116f4-ac53-4a41-ad68-6606788e674c	0	appliances	environment
4	c6946211-ce70-4228-95ce-fd7fccdde63c	0	construction	tools

	cat_2	cat_3	timestamp	ts_hour	ts_minute	ts_weekday	ts_day	\
0	<NA>	<NA>	2019-12-01 00:00:28	0	0	6	1	
1	vacuum	<NA>	2019-12-01 00:00:39	0	0	6	1	
2	vacuum	<NA>	2019-12-01 00:00:40	0	0	6	1	
3	vacuum	<NA>	2019-12-01 00:00:41	0	0	6	1	
4	light	<NA>	2019-12-01 00:01:56	0	1	6	1	

	ts_month	ts_year
0	12	2019
1	12	2019
2	12	2019
3	12	2019
4	12	2019

```
[19]: ###ToDo
def explore_cat(df, cats):
    df_agg = df_train[cats + ['target']].groupby(cats).agg(['mean', 'count']).
    ↪reset_index()
    df_agg.columns = cats + ['mean', 'count']
    print(df_agg.sort_values('count', ascending=False).head(20))

cats = ['product_id', 'user_id']
explore_cat(df_train, cats)
```

	product_id	user_id	mean	count
640663	1004767	545442548	0.000000	807
620114	1004767	525325337	0.000000	753
1560100	1005107	553431815	0.599185	736
4882910	15300303	512875426	0.000000	709
2021336	1005174	563599039	0.931238	509
1148644	1004873	515032042	0.002041	490
3626024	4804718	536911254	0.000000	471
4154253	8800045	557590749	0.000000	380
3076839	3601537	578263741	0.000000	363
839819	1004833	564068124	0.793872	359
1477094	1005100	611998200	0.000000	333
6442	1002524	515598234	0.677215	316
348093	1004249	513901034	0.648562	313
1301035	1005008	521558076	0.003236	309
1941657	1005161	512924342	0.537162	296
1593381	1005115	516010934	0.750000	288

3631591	4804718	576154686	0.550523	287
5935044	100007950	515481166	0.000000	287
3062423	3601489	513824664	0.000000	275
85419	1002544	545376441	0.896296	270

```
[20]: ##### Solution #####
```

```
[23]: ##### Solution End #####
```

2.4 Optimization

There is not much optimization technique to apply. We will “chain” the idea of combining categories with other Feature Engineering techniques, which does NOT require us to actually combine and store the new feature in the dataset. Instead, we will create features based on the combined categories directly and won’t store the combined categories as a separate feature. One advice is to use cuDF instead of pandas. Analyzing the dataset requires calculating different groupby combination multiple times by a data scientist. GPU acceleration can significantly speed-up the calculations and enables you to run more comparisons.

```
[24]: big_df = df_train.to_pandas()
      big_data = df_train
```

```
[25]: print('Pandas Shape:' + str(big_df.shape))
      print('cudf Shape:' + str(big_df.shape))
```

```
Pandas Shape:(11461357, 19)
cudf Shape:(11461357, 19)
```

```
[26]: %%time

big_df.groupby(['cat_0', 'cat_1', 'cat_2', 'cat_3', 'brand']).target.
      ↪agg(['mean', 'count'])
print('')
```

```
CPU times: user 3.02 s, sys: 616 ms, total: 3.63 s
Wall time: 3.63 s
```

```
[27]: %%time

big_data.groupby(['cat_0', 'cat_1', 'cat_2', 'cat_3', 'brand']).target.
      ↪agg(['mean', 'count'])
print('')
```

```
CPU times: user 812 ms, sys: 800 ms, total: 1.61 s
Wall time: 1.61 s
```


A dataset with 12M rows is ~4-6x faster on GPU with cuDF as on CPU with pandas. This difference can even increase with larger dataset size as the groupby operation is not linear in complexity.

We shutdown the kernel.

```
[28]: app = IPython.Application.instance()  
      app.kernel.do_shutdown(False)
```

```
[28]: {'status': 'ok', 'restart': False}
```