

## 03\_2\_Categorify

January 28, 2021

```
[ ]: # The MIT License (MIT)

# Copyright (c) 2020, NVIDIA CORPORATION.

# Permission is hereby granted, free of charge, to any person obtaining a copy
# of
# this software and associated documentation files (the "Software"), to deal in
# the Software without restriction, including without limitation the rights to
# use, copy, modify, merge, publish, distribute, sublicense, and/or sell copies
# of
# the Software, and to permit persons to whom the Software is furnished to do
# so,
# subject to the following conditions:

# The above copyright notice and this permission notice shall be included in all
# copies or substantial portions of the Software.

# THE SOFTWARE IS PROVIDED "AS IS", WITHOUT WARRANTY OF ANY KIND, EXPRESS OR
# IMPLIED, INCLUDING BUT NOT LIMITED TO THE WARRANTIES OF MERCHANTABILITY,
# FITNESS
# FOR A PARTICULAR PURPOSE AND NONINFRINGEMENT. IN NO EVENT SHALL THE AUTHORS OR
# COPYRIGHT HOLDERS BE LIABLE FOR ANY CLAIM, DAMAGES OR OTHER LIABILITY, WHETHER
# IN AN ACTION OF CONTRACT, TORT OR OTHERWISE, ARISING FROM, OUT OF OR IN
# CONNECTION WITH THE SOFTWARE OR THE USE OR OTHER DEALINGS IN THE SOFTWARE
```

### 1 Tutorial: Feature Engineering for Recommender Systems

### 2 3. Feature Engineering - Categorical

#### 2.1 3.2. Categorify

```
[1]: import IPython

import pandas as pd
import cudf
import numpy as np
```

```
import cupy
import matplotlib.pyplot as plt

df_train = cudf.read_parquet('./data/train.parquet')
df_valid = cudf.read_parquet('./data/valid.parquet')
df_test = cudf.read_parquet('./data/test.parquet')
```

```
[2]: df_train.head()
```

```
[2]:
```

		event_time	event_type	product_id	brand	price	user_id	\
0	2019-12-01	00:00:28 UTC	cart	17800342	zeta	66.90	550465671	
1	2019-12-01	00:00:39 UTC	cart	3701309	polaris	89.32	543733099	
2	2019-12-01	00:00:40 UTC	cart	3701309	polaris	89.32	543733099	
3	2019-12-01	00:00:41 UTC	cart	3701309	polaris	89.32	543733099	
4	2019-12-01	00:01:56 UTC	cart	1004767	samsung	235.60	579970209	

  

		user_session	target	cat_0	cat_1	\
0	22650a62-2d9c-4151-9f41-2674ec6d32d5	0	computers	desktop		
1	a65116f4-ac53-4a41-ad68-6606788e674c	0	appliances	environment		
2	a65116f4-ac53-4a41-ad68-6606788e674c	0	appliances	environment		
3	a65116f4-ac53-4a41-ad68-6606788e674c	0	appliances	environment		
4	c6946211-ce70-4228-95ce-fd7fccdde63c	0	construction	tools		

  

	cat_2	cat_3	timestamp	ts_hour	ts_minute	ts_weekday	ts_day	\
0	<NA>	<NA>	2019-12-01 00:00:28	0	0	6	1	
1	vacuum	<NA>	2019-12-01 00:00:39	0	0	6	1	
2	vacuum	<NA>	2019-12-01 00:00:40	0	0	6	1	
3	vacuum	<NA>	2019-12-01 00:00:41	0	0	6	1	
4	light	<NA>	2019-12-01 00:01:56	0	1	6	1	

  

	ts_month	ts_year
0	12	2019
1	12	2019
2	12	2019
3	12	2019
4	12	2019

```
[3]: cat = 'product_id'
```

## 2.2 Theory

*Categorifying* is required for using categorical features in deep learning models with Embedding layers. An Embedding layer encodes the category into a hidden latent vector with a smaller dimension. Categorical features can be from datatype String or Integer. The Embedding layer requires that categorical features are continuous, positive Integers from 0 to  $|C|$  (number of unique category values).

There are 164453 unique product values but the ProductIDs range from 1000894 to 100144608.

```
[4]: df_train[cat].unique()
```

```
[4]: 0          1000894
      1          1000978
      2          1001588
      3          1001605
      4          1001606
      ...
      164448      100143856
      164449      100143867
      164450      100144046
      164451      100144443
      164452      100144608
      Name: product_id, Length: 164453, dtype: int64
```

Using factorize creates continuous Integers from a categorical feature.

```
[5]: codes, uniques = df_train[cat].factorize()
```

```
[6]: codes
```

```
[6]: 0          65426
      1          10158
      2          10158
      3          10158
      4           775
      ...
      11461352      862
      11461353     1064
      11461354       775
      11461355     10158
      11461356     91487
      Length: 11461357, dtype: int32
```

```
[7]: codes.unique()
```

```
[7]: 0          0
      1          1
      2          2
      3          3
      4          4
      ...
      164448     164448
      164449     164449
      164450     164450
      164451     164451
      164452     164452
```

Length: 164453, dtype: int32

Another important reason to Categorify categorical features is to reduce the size of the dataset. Often categorical features are of the datatype String and sometimes, they are hashed to protect the user / dataset privacy.

```
[8]: import hashlib
     from sys import getsizeof
```

For example, we can hash the Integer 0 to a md5 hash

```
[9]: hashlib.md5(b'0').hexdigest()
```

```
[9]: 'cfcd208495d565ef66e7dff9f98764da'
```

We can hash the full product\_id column

```
[10]: hashSeries = df_train[cat].to_pandas().apply(lambda x: hashlib.
        ↳md5(bytes(str(x), encoding='utf-8')).hexdigest())
```

```
[11]: hashSeries
```

```
[11]: 0          5ebc4b45850c48658af86229318ccbea
      1          4b9dde859aa2809cc367fc44aa05eb4a
      2          4b9dde859aa2809cc367fc44aa05eb4a
      3          4b9dde859aa2809cc367fc44aa05eb4a
      4          e5e26a76d8aee9c4f8b4cd9cb8633577
      ...
      11461352    4202ee67e0c3f9b1ebcfb622d9974e07
      11461353    5a06406ed78aab3c94bfefcdeb528eaf
      11461354    e5e26a76d8aee9c4f8b4cd9cb8633577
      11461355    4b9dde859aa2809cc367fc44aa05eb4a
      11461356    388318441a4807e254acf9c3f207969d
      Name: product_id, Length: 11461357, dtype: object
```

```
[12]: getsizeof(hashSeries)
```

```
[12]: 1020060933
```

```
[13]: codes, uniques = hashSeries.factorize()
```

```
[14]: getsizeof(pd.DataFrame(codes)[0])
```

```
[14]: 91691016
```

We require only 9% of the original DataSeries memory.

```
[15]: 91691016/1020060933
```

```
[15]: 0.08988778320363339
```

Finally, we can prevent overfitting for low frequency categories. Categories with low frequency can be grouped together to an new category called 'other'. In the previous exercise we learned that it is powerful to combine categorical features together to create a new feature. However, combining categories increases the cardinality of the new feature and the number of observations per category will decrease. Therefore, we can apply a threshold to group all categories with lower frequency count to the new category. In addition, categories, which occur in the validation dataset and do not occur in the training dataset, should be mapped to the 'other' category as well. We use in our example the categoryIds 0 or 1 for a placeholder for the low frequency and unknown category. Then our function is independent of the cardinality of the categorical feature and we do not keep records of the cardinality to know the low frequency/unknown category.

In our dataset, we see that multiple product\_ids occur only once in the training dataset. Our model would overfit to these low frequent categories.

```
[16]: df_train[cat].value_counts()
```

```
[16]: 1004767      317711
      1005115      251189
      1004856      227432
      4804056      224545
      1005100      180072
      ...
      100143590         1
      100143856         1
      100143867         1
      100144046         1
      100144443         1
      Name: product_id, Length: 164453, dtype: int32
```

```
[17]: freq = df_train[cat].value_counts()
```

```
[18]: freq = freq.reset_index()
      freq.columns = [cat, 'count']
      freq = freq.reset_index()
      freq.columns = [cat + '_Categorify', cat, 'count']
      freq_filtered = freq[freq['count'] > 5]
      freq_filtered[cat + '_Categorify'] = freq_filtered[cat + '_Categorify'] + 1
      freq_filtered = freq_filtered.drop('count', axis=1)
      df_train = df_train.merge(freq_filtered, how='left', on=cat)
      df_train[cat + '_Categorify'] = df_train[cat + '_Categorify'].fillna(0)
```

```
[19]: df_train['product_id_Categorify'].min(), df_train['product_id_Categorify'].
      ↪max(), df_train['product_id_Categorify'].drop_duplicates().shape
```

```
[19]: (0, 76404, (76405,))
```

We need to apply the categorify to our validation and test sets.

```
[20]: df_valid = df_valid.merge(freq_filtered, how='left', on=cat)
      df_valid[cat + '_Categorify'] = df_valid[cat + '_Categorify'].fillna(0)

      df_test = df_test.merge(freq_filtered, how='left', on=cat)
      df_test[cat + '_Categorify'] = df_test[cat + '_Categorify'].fillna(0)
```

**Summary** Categorify is important to enable deep learning models to use categorical features

Categorify can significantly reduce the dataset size by transforming categorical features from String datatypes to Integer datatypes

Categorify can prevent overfitting by grouping categories with low frequency into one category together

## 2.3 Practice

Now, it is your turn

### ToDo:

Categorify the category features brand

Apply a frequency treshhold of minimum 20

Map low frequency categories to the id=0

Map unkown categories to the id=1 in the validation and test set

### Question:

How many data points have an unknown category in the test dataset?

How many data points have a low frequency category in the test dataset?

How many data points have a low frequency category in the training dataset?

```
[21]: ### ToDo
```

```
[22]: ##### Solution #####
```

```
[27]: ##### Solution End #####
```

## 2.4 Optimization

Let's compare the runtime between pandas and cuDF. The implementation depends only on the DataFrame object (calling function of the object) and does not require any pd / cuDF function. Therefore, we can use the same implementation and just use pandas.DataFrame and cuDF.DataFrame.

```
[28]:
```

```
def categorify(df_train, df_valid, df_test, cat, freq_treshhold=20,
    ↪unkown_id=1, lowfrequency_id=0):
    freq = df_train[cat].value_counts()
    freq = freq.reset_index()
    freq.columns = [cat, 'count']
    freq = freq.reset_index()
    freq.columns = [cat + '_Categorify', cat, 'count']
    freq[cat + '_Categorify'] = freq[cat + '_Categorify']+2
    freq.loc[freq['count']<freq_treshhold, cat + '_Categorify'] =
    ↪lowfrequency_id

    freq = freq.drop('count', axis=1)
    df_train = df_train.merge(freq, how='left', on=cat)
    df_train[cat + '_Categorify'] = df_train[cat + '_Categorify'].
    ↪fillna(unkown_id)

    df_valid = df_valid.merge(freq, how='left', on=cat)
    df_valid[cat + '_Categorify'] = df_valid[cat + '_Categorify'].
    ↪fillna(unkown_id)

    df_test = df_test.merge(freq, how='left', on=cat)
    df_test[cat + '_Categorify'] = df_test[cat + '_Categorify'].
    ↪fillna(unkown_id)
```

```
[29]: df_train_pd = df_train.to_pandas()
df_valid_pd = df_valid.to_pandas()
df_test_pd = df_test.to_pandas()
```

```
[30]: %%time

categorify(df_train_pd, df_valid_pd, df_test_pd, 'user_id')
```

CPU times: user 15.5 s, sys: 5.77 s, total: 21.3 s  
Wall time: 21.2 s

```
[31]: %%time

categorify(df_train, df_valid, df_test, 'user_id')
```

CPU times: user 168 ms, sys: 296 ms, total: 464 ms  
Wall time: 463 ms

In our experiments, running the same implementation is 63x times faster with cuDF instead of pandas.

We shutdown the kernel.

```
[32]: app = IPython.Application.instance()  
      app.kernel.do_shutdown(False)
```

```
[32]: {'status': 'ok', 'restart': False}
```