

## 05\_2\_TimeSeries\_Differences

January 28, 2021

```
[ ]: # The MIT License (MIT)

# Copyright (c) 2020, NVIDIA CORPORATION.

# Permission is hereby granted, free of charge, to any person obtaining a copy
# of
# this software and associated documentation files (the "Software"), to deal in
# the Software without restriction, including without limitation the rights to
# use, copy, modify, merge, publish, distribute, sublicense, and/or sell copies
# of
# the Software, and to permit persons to whom the Software is furnished to do
# so,
# subject to the following conditions:

# The above copyright notice and this permission notice shall be included in all
# copies or substantial portions of the Software.

# THE SOFTWARE IS PROVIDED "AS IS", WITHOUT WARRANTY OF ANY KIND, EXPRESS OR
# IMPLIED, INCLUDING BUT NOT LIMITED TO THE WARRANTIES OF MERCHANTABILITY,
# FITNESS
# FOR A PARTICULAR PURPOSE AND NONINFRINGEMENT. IN NO EVENT SHALL THE AUTHORS OR
# COPYRIGHT HOLDERS BE LIABLE FOR ANY CLAIM, DAMAGES OR OTHER LIABILITY, WHETHER
# IN AN ACTION OF CONTRACT, TORT OR OTHERWISE, ARISING FROM, OUT OF OR IN
# CONNECTION WITH THE SOFTWARE OR THE USE OR OTHER DEALINGS IN THE SOFTWARE
```

## 1 Tutorial: Feature Engineering for Recommender Systems

### 2 5. Feature Engineering

#### 2.1 5.2. Differences

```
[1]: import IPython

import pandas as pd
import numpy as np

import cudf
```

```
import cupy
```

```
np.random.seed(42)
```

```
[2]: itemid = [1000001]*10 + [1000002]*5 + [1000001]*5 + [1000002]*5 + [1000001]*1 +  
    ↪ [1000002]*1 + [1000001]*2 + [1000002]*2  
    itemid += [1000001]*3 + [1000002]*2 + [1000001]*1 + [1000002]*1 + [1000001]*6 +  
    ↪ [1000002]*3 + [1000001]*2 + [1000002]*2  
    userid = np.random.choice(list(range(10000)), len(itemid))  
    action = np.random.choice(list(range(2)), len(itemid), p=[0.2, 0.8])  
  
    price = [100.00]*10 + [25.00]*5 + [100.00]*5 + [30.00]*5 + [125.00]*1 + [30.  
    ↪ 00]*1 + [125.00]*2 + [30.00]*2  
    price += [110.00]*3 + [30.00]*2 + [110.00]*1 + [20.00]*1 + [90.00]*6 + [20.  
    ↪ 00]*3 + [90.00]*2 + [20.00]*2  
  
    timestamp = [pd.to_datetime('2020-01-01')]*15  
    timestamp += [pd.to_datetime('2020-01-02')]*10  
    timestamp += [pd.to_datetime('2020-01-03')]*2  
    timestamp += [pd.to_datetime('2020-01-04')]*4  
    timestamp += [pd.to_datetime('2020-01-05')]*5  
    timestamp += [pd.to_datetime('2020-01-07')]*2  
    timestamp += [pd.to_datetime('2020-01-08')]*9  
    timestamp += [pd.to_datetime('2020-01-09')]*4  
  
    data = pd.DataFrame({  
        'itemid': itemid,  
        'userid': userid,  
        'price': price,  
        'action': action,  
        'timestamp': timestamp  
    })  
  
    data = cudf.from_pandas(data)
```

## 2.2 Theory

Another category of powerful features is to calculate the differences to previous datapoints based on a timestamp. For example, we can calculate if the price changed of a product and how much the price change was.

```
[3]: data[data['itemid']==1000001].head(10)
```

```
[3]:
```

	itemid	userid	price	action	timestamp
0	1000001	7270	100.0	1	2020-01-01
1	1000001	860	100.0	1	2020-01-01
2	1000001	5390	100.0	0	2020-01-01

3	1000001	5191	100.0	1	2020-01-01
4	1000001	5734	100.0	0	2020-01-01
5	1000001	6265	100.0	1	2020-01-01
6	1000001	466	100.0	1	2020-01-01
7	1000001	4426	100.0	1	2020-01-01
8	1000001	5578	100.0	1	2020-01-01
9	1000001	8322	100.0	0	2020-01-01

Tree-based or deep learning based models have difficulties processing these relationships on their own. Providing the models with these features can significantly improve the performance.

```
[4]: offset = 1
data_shift = data[['itemid', 'timestamp', 'price']].groupby(['itemid',
↳ 'timestamp']).mean().reset_index()
data_shift.columns = ['itemid', 'timestamp', 'mean']
data_shift['mean_' + str(offset)] = data_shift['mean'].shift(1)
data_shift.loc[data_shift['itemid']!=data_shift['itemid'].shift(1), 'mean_' +
↳ str(offset)] = None
data_shift['diff_' + str(offset)] = data_shift['mean'] - data_shift['mean_' +
↳ str(offset)]
```

```
[5]: data_shift.head(10)
```

```
[5]:
```

	itemid	timestamp	mean	mean_1	diff_1
0	1000001	2020-01-01	100.0	<NA>	<NA>
1	1000001	2020-01-02	100.0	100.0	0.0
2	1000001	2020-01-03	125.0	100.0	25.0
3	1000001	2020-01-04	125.0	125.0	0.0
4	1000001	2020-01-05	110.0	125.0	-15.0
5	1000001	2020-01-07	110.0	110.0	0.0
6	1000001	2020-01-08	90.0	110.0	-20.0
7	1000001	2020-01-09	90.0	90.0	0.0
8	1000002	2020-01-01	25.0	<NA>	<NA>
9	1000002	2020-01-02	30.0	25.0	5.0

```
[6]: data_shift.columns = ['itemid', 'timestamp', 'c1', 'c2', 'price_diff_1']
data_shift.drop(['c1', 'c2'], inplace=True).head(10)
```

```
[6]:
```

	itemid	timestamp	price_diff_1
0	1000001	2020-01-01	<NA>
1	1000001	2020-01-02	0.0
2	1000001	2020-01-03	25.0
3	1000001	2020-01-04	0.0
4	1000001	2020-01-05	-15.0
5	1000001	2020-01-07	0.0
6	1000001	2020-01-08	-20.0
7	1000001	2020-01-09	0.0

```

8  1000002  2020-01-01      <NA>
9  1000002  2020-01-02       5.0

```

```
[7]: data = data.merge(data_shift, how='left', on=['itemid', 'timestamp'])
```

```
[8]: data.head()
```

```

[8]:   itemid  userid  price  action  timestamp  price_diff_1
0  1000001    4658   110.0      0  2020-01-05          -15.0
1  1000001    1899   110.0      0  2020-01-05          -15.0
2  1000002    7734    30.0      1  2020-01-05           0.0
3  1000002    1267    30.0      1  2020-01-05           0.0
4  1000001    1528   110.0      1  2020-01-07           0.0

```

We can combine techniques of TimeSeries data and chain them together. For example, we can calculate the # of purchases per item and then compare the previous week with a the week, 2, 3 or 5 weeks ago. We can recognize patterns over time.

## 2.3 Practise

```

[9]: import pandas as pd
import cudf
import numpy as np
import cupy
import matplotlib.pyplot as plt

df_train = cudf.read_parquet('./data/train.parquet')
df_valid = cudf.read_parquet('./data/valid.parquet')
df_test = cudf.read_parquet('./data/test.parquet')

df_train['brand'] = df_train['brand'].fillna('UNKNOWN')
df_valid['brand'] = df_valid['brand'].fillna('UNKNOWN')
df_test['brand'] = df_test['brand'].fillna('UNKNOWN')

df_train['cat_0'] = df_train['cat_0'].fillna('UNKNOWN')
df_valid['cat_0'] = df_valid['cat_0'].fillna('UNKNOWN')
df_test['cat_0'] = df_test['cat_0'].fillna('UNKNOWN')

df_train['cat_1'] = df_train['cat_1'].fillna('UNKNOWN')
df_valid['cat_1'] = df_valid['cat_1'].fillna('UNKNOWN')
df_test['cat_1'] = df_test['cat_1'].fillna('UNKNOWN')

df_train['cat_2'] = df_train['cat_2'].fillna('UNKNOWN')
df_valid['cat_2'] = df_valid['cat_2'].fillna('UNKNOWN')
df_test['cat_2'] = df_test['cat_2'].fillna('UNKNOWN')

```

cuDF does not support date32, right now. We use pandas to transform the timestamp in only date values.

```
[10]: df_train['date'] = cudf.from_pandas(pd.to_datetime(df_train['timestamp']).
      ↪to_pandas()).dt.date)
```

```
/conda/envs/nvtabular/lib/python3.7/site-
packages/cudf/core/column/column.py:1396: UserWarning: Date32 values are not yet
supported so this will be typecast to a Date64 value
  UserWarning,
```

**ToDo:**

Let's get the price difference of the previous price to the current price per item

## 2.4 Optimisation

Let's compare a CPU with the GPU version.

```
[14]: def difference_feature(df, offset):
      data_shift = df[['product_id', 'date', 'price']].groupby(['product_id',
      ↪'date']).mean().reset_index()
      data_shift.columns = ['product_id', 'date', 'mean']
      data_shift['mean_' + str(offset)] = data_shift['mean'].shift(offset)
      data_shift.loc[data_shift['product_id']!=data_shift['product_id'].
      ↪shift(offset), 'mean_' + str(offset)] = None
      data_shift['diff_' + str(offset)] = data_shift['mean'] - data_shift['mean_' +
      ↪str(offset)]
      data_shift.columns = ['product_id', 'date', 'c1', 'c2', 'price_diff_' +
      ↪str(offset)]
      data_shift.drop(['c1', 'c2'], axis=1, inplace=True)
      df = df.merge(data_shift, how='left', on=['product_id', 'date'])
```

```
[15]: df_train_pd = df_train.to_pandas()
```

```
[16]: %%time

_ = difference_feature(df_train_pd, 1)
```

```
CPU times: user 10.2 s, sys: 4.81 s, total: 15 s
Wall time: 15 s
```

```
[17]: %%time

_ = difference_feature(df_train, 1)
```

```
CPU times: user 196 ms, sys: 252 ms, total: 448 ms
Wall time: 444 ms
```

In our experiments, we achieved a speedup of 43.1s

We shutdown the kernel.

```
[1]: import IPython

app = IPython.Application.instance()
app.kernel.do_shutdown(False)
```

```
[1]: {'status': 'ok', 'restart': False}
```