

06__1__Intro__Dask

January 28, 2021

```
[ ]: # The MIT License (MIT)

# Copyright (c) 2020, NVIDIA CORPORATION.

# Permission is hereby granted, free of charge, to any person obtaining a copy
# of
# this software and associated documentation files (the "Software"), to deal in
# the Software without restriction, including without limitation the rights to
# use, copy, modify, merge, publish, distribute, sublicense, and/or sell copies
# of
# the Software, and to permit persons to whom the Software is furnished to do
# so,
# subject to the following conditions:

# The above copyright notice and this permission notice shall be included in all
# copies or substantial portions of the Software.

# THE SOFTWARE IS PROVIDED "AS IS", WITHOUT WARRANTY OF ANY KIND, EXPRESS OR
# IMPLIED, INCLUDING BUT NOT LIMITED TO THE WARRANTIES OF MERCHANTABILITY,
# FITNESS
# FOR A PARTICULAR PURPOSE AND NONINFRINGEMENT. IN NO EVENT SHALL THE AUTHORS OR
# COPYRIGHT HOLDERS BE LIABLE FOR ANY CLAIM, DAMAGES OR OTHER LIABILITY, WHETHER
# IN AN ACTION OF CONTRACT, TORT OR OTHERWISE, ARISING FROM, OUT OF OR IN
# CONNECTION WITH THE SOFTWARE OR THE USE OR OTHER DEALINGS IN THE SOFTWARE
```

1 Tutorial: Feature Engineering for Recommender Systems

2 6. Scaling to Production Systems

2.1 6.1. Introduction to dask and dask_cudf

2.2 Theory

Acknowledgement: Much of the introductory material included here is borrowed from other Dask documentation and tutorials. - [“Dask Video Tutorial” - YouTube link](#) - [Introduction To Dask by Richard \(Rick\) Zamora](#)

Other useful Dask resources: - [Dask.org](#) - [Tutorial pages](#) - [GitHub Tutorial](#)

2.3 What is Dask

Very Short Answer: Dask is an open-source library designed to natively scale Python code.

Slightly-Longer Short Answer: Dask is a task-based library for parallel scheduling and execution. Although it is certainly possible to use the task-scheduling machinery directly to implement customized parallel workflows (we do it in NVTabular), most users only interact with Dask through a *Dask Collection API*. The most popular “collection” API’s include:

- **Dask DataFrame:** Dask-based version of the [Pandas](#) DataFrame/Series API. Note that `dask_cudf` is just a wrapper around this collection module (`dask.dataframe`).
- **Dask Array:** Dask-based version of the NumPy array API
- **Dask Bag:** *Similar to* a Dask-based version of PyToolz or a Pythonic version of PySpark RDD

For example, Dask DataFrame provides a convenient API for decomposing large pandas (or cuDF) DataFrame/Series objects into a collection of DataFrame *partitions*. This tutorial will focus mostly on this particular Dask collection (since it is the basis for `dask_cudf`). However, instead of relying only on the established `dask.dataframe` API, we will also see how it is possible (perhaps easy) to implement a custom task graph to operate on Dask-DataFrame objects when necessary.

2.3.1 Dask Uses DAGs Internally

Before we start writing any code, it is useful to understand (on a basic level) how Dask actually works. When an application or library uses a Dask collection API (like Dask DataFrame), they are typically using that API to construct a directed acyclic graph (DAG) of tasks. Once a DAG is constructed, the **core** Dask API can be used (either directly or implicitly through the collection API) to schedule and execute the DAG on one or more threads/processes.

In other words, Dask provides various APIs to:

1. Construct a DAG of “tasks”
2. Schedule/execute those DAGs
3. (Optionally) Spin up a dedicated worker and scheduler processes to enable distributed execution

2.3.2 Important Components of the “Dask Ecosystem”

The components of the Dask ecosystem that are most critical for NVTabular (and will be discussed in this tutorial) are:

- **dask** (core Dask library): [[GitHub Repo](#)] This is the core Dask library. It also contains the Dask Dataframe API (`dask.dataframe`)
- **dask_cudf**: [[GitHub Repo](#)] This is effectively a wrapper around the `dask.dataframe` module defined in the core Dask library. Note that a `dask_cudf.DataFrame` object should be thought of as a `dask.dataframe.DataFrame` object, but with the underlying partitions being `cudf.DataFrame`’s (rather than `pandas.DataFrame`)
- **distributed**: [[GitHub Repo](#)] Distributed version of the Dask execution model (includes the necessary code for scheduling, execution and communication between distributed processes). This library does not deal with the construction of DAGs, just with the scheduling and execution of DAGs on distributed *workers*.

- `dask_cuda`: [\[GitHub Repo\]](#) Provides various utilities to improve deployment and management of distributed Dask *workers* on CUDA-enabled systems.

2.4 HandsOn

Before we get started, it is convenient to create a simple `dask.distributed` client. If we work with a small dataset, then it is not necessary to initialize a `dask.distributed` client. The code should run in the same way.

```
[1]: import dask
      from dask.distributed import Client, LocalCluster
      import dask.dataframe as dd
```

```
[2]: client = Client(n_workers=8,
                    threads_per_worker=1,
                    memory_limit='50GB',
                    ip='127.0.0.1')

client
```

```
[2]: <Client: 'tcp://127.0.0.1:40757' processes=8 threads=8, memory=132.13 GB>
```

```
[3]: %%time

ddf_train = dd.read_parquet('./data/train.parquet', blocksize=12e3)
ddf_valid = dd.read_parquet('./data/valid.parquet', blocksize=12e3)
```

```
CPU times: user 28 ms, sys: 0 ns, total: 28 ms
Wall time: 27.4 ms
```

```
[4]: ddf_train
```

```
[4]: Dask DataFrame Structure:
              event_time event_type product_id  brand  price user_id
user_session target  cat_0  cat_1  cat_2  cat_3 timestamp ts_hour ts_minute
ts_weekday ts_day ts_month ts_year
npartitions=1
              object      object      int64  object  float64  int64
object  int64  object  object  object  object  object  int64  int64
int64  int64  int64  int64
...
...  ...  ...  ...  ...  ...  ...
...  ...  ...  ...
Dask Name: read-parquet, 1 tasks
```

Here we have created a `dask.dataframe.DataFrame` object called `ddf_train` and `ddf_valid`. Both are essentially a (**lazy**) collection of pandas dataframes. Dask loaded the metadata (DataFrame schema) but did not load any data in-memory. Each pandas dataframe in this collection is called a **partition**. We can access this property (the total number of partitions) using the

DataFrame.npartitions attribute.

It is absolutely critical to recognize that `ddf_train` and `ddf_valid` are *not* actually backed by *in-memory* pandas data, but instead by a DAG of tasks. This DAG (accessible via `ddf.dask`) specifies the exact network of operations needed to produce the underlying partitions.

```
[5]: ddf_train._meta
```

```
[5]: Empty DataFrame
Columns: [event_time, event_type, product_id, brand, price, user_id,
user_session, target, cat_0, cat_1, cat_2, cat_3, timestamp, ts_hour, ts_minute,
ts_weekday, ts_day, ts_month, ts_year]
Index: []
```

Let's work on some examples: Simplified Target Encoding 1. We combine two columns `cat_2` and `brand` 2. We TargetEncode the new column `cat_2_brand` 3. We merge the counts back to the train and validation dataset 4. We overwrite counts with less than 20 for `cat_2_brand` with `global_mean`

We can see that the execution time is 117ms - meaning that `dask` has registered the operations but hasn't executed them.

```
[6]: %%time

ddf_train['cat_2_brand'] = ddf_train['cat_2'].astype(str) + '_' +
    ↳ddf_train['brand'].astype(str)
ddf_valid['cat_2_brand'] = ddf_valid['cat_2'].astype(str) + '_' +
    ↳ddf_valid['brand'].astype(str)

ddf_train_group = ddf_train[['cat_2_brand', 'target']].groupby(['cat_2_brand']).
    ↳agg(['count', 'mean'])
ddf_train_group = ddf_train_group.reset_index()
ddf_train_group.columns = ['cat_2_brand', 'TE_count', 'TE_mean']
ddf_train = ddf_train.merge(ddf_train_group, how='left', on='cat_2_brand')
ddf_valid = ddf_valid.merge(ddf_train_group, how='left', on='cat_2_brand')
global_mean = ddf_train['target'].mean()
ddf_train['TE_mean'] = ddf_train.TE_mean.where(ddf_train['TE_count']>20,
    ↳global_mean)
ddf_valid['TE_mean'] = ddf_valid.TE_mean.where(ddf_valid['TE_count']>20,
    ↳global_mean)
```

CPU times: user 60 ms, sys: 0 ns, total: 60 ms

Wall time: 56.7 ms

We can compute the task graph by calling `.compute()` or `.persist()`

```
[7]: %%time

ddf_train.compute()
```

```
ddf_valid.compute()
```

CPU times: user 15.1 s, sys: 9.68 s, total: 24.8 s

Wall time: 3min 14s

```
[7]:
```

		event_time	event_type	product_id	brand	price	\
0	2020-03-01	00:00:59 UTC	cart	6902464	zlatek	49.91	
1	2020-03-01	00:01:20 UTC	cart	1002544	apple	397.10	
2	2020-03-01	00:01:52 UTC	cart	1003316	apple	823.70	
3	2020-03-01	00:02:14 UTC	cart	16600067	rivertoys	422.15	
4	2020-03-01	00:02:15 UTC	cart	3701428	arnica	69.24	
...
2461714	2020-03-31	23:57:47 UTC	purchase	24100293	cocochocho	2.65	
2461715	2020-03-31	23:58:19 UTC	purchase	100049773	None	234.96	
2461716	2020-03-31	23:58:20 UTC	purchase	3700689	samsung	223.92	
2461717	2020-03-31	23:59:19 UTC	purchase	100077607	vitek	100.36	
2461718	2020-03-31	23:59:27 UTC	purchase	100068493	samsung	319.41	

	user_id	user_session	target	\
0	531574188	48714293-b3f9-4946-8135-eb1ea05ead74	0	
1	622090790	fb5b918c-f1f6-48d9-bcf4-7eb46e83fc6b	0	
2	622090543	b821ee79-96fe-4979-be9d-21ee2e6777c3	0	
3	616437533	aad023bc-c858-47ab-a3a7-ff4654f11b9a	0	
4	516454226	ee22b80c-ed3e-3c83-d397-fb69a44d4864	0	
...
2461714	513094047	d27f822c-f707-4956-a6c3-4ad8fec00cc7	1	
2461715	620580925	c33fde42-a5de-4a1f-9e1c-2ac7518a7d41	1	
2461716	514905289	e40783c5-7b21-429f-99af-539d2842e6d3	1	
2461717	633281427	667a8535-221c-4169-aab4-a1972610f102	1	
2461718	635165435	861f2378-076f-4ddd-85e3-9844923d03a9	1	

	cat_0	cat_1	...	timestamp	ts_hour	\
0	electronics	telephone	...	2020-03-01 00:00:59	0	
1	construction	tools	...	2020-03-01 00:01:20	0	
2	construction	tools	...	2020-03-01 00:01:52	0	
3	sport	trainer	...	2020-03-01 00:02:14	0	
4	appliances	environment	...	2020-03-01 00:02:15	0	
...
2461714	appliances	personal	...	2020-03-31 23:57:47	23	
2461715	None	None	...	2020-03-31 23:58:19	23	
2461716	appliances	environment	...	2020-03-31 23:58:20	23	
2461717	appliances	environment	...	2020-03-31 23:59:19	23	
2461718	construction	tools	...	2020-03-31 23:59:27	23	

	ts_minute	ts_weekday	ts_day	ts_month	ts_year	cat_2_brand	\
0	0	6	1	3	2020	nan_zlatek	
1	1	6	1	3	2020	light_apple	

2	1	6	1	3	2020	light_apple
3	2	6	1	3	2020	nan_rivertoys
4	2	6	1	3	2020	vacuum_arnica
...
2461714	57	1	31	3	2020	massager_cocochoco
2461715	58	1	31	3	2020	nan_nan
2461716	58	1	31	3	2020	vacuum_samsung
2461717	59	1	31	3	2020	vacuum_vitek
2461718	59	1	31	3	2020	light_samsung

	TE_count	TE_mean
0	607.0	0.258649
1	1013391.0	0.469441
2	1013391.0	0.469441
3	10564.0	0.104411
4	4450.0	0.325393
...
2461714	82.0	0.146341
2461715	521515.0	0.277106
2461716	68239.0	0.392459
2461717	11950.0	0.340000
2461718	1212393.0	0.481047

[2461719 rows x 22 columns]

```
[8]: client.close()
```

About compute: The `compute` method is [defined for all Dask collections](#). For Dask DataFrame, this method will (1) trigger the execution of the graph and (2) convert the Dask DataFrame into a **single** Pandas DataFrame. *This means that you should be sure the pandas equivalent of `ddf` will fit in memory before you use `compute`!*

Using persist

Since the `compute` method will convert your Dask DataFrame to a Pandas DataFrame, it is typically a **bad** idea to use `compute` on larger-than-memory (LTM) datasets. In NVTabular, we do use a `compute` method, but never on a full Dask/dask_cudf DataFrame object. Instead, we use `compute` to trigger the collection/reduction of an aggregated statistics dictionary, and/or to write out a processed dataset.

In order to execute the `ddf` DAG **without** converting it to a single pandas DataFrame, you need to use the [persist method](#). This function is particularly useful when using distributed systems, because the results will be kept in distributed memory, rather than returned to the local process as with `compute`. It will also allow the distributed cluster to clean up data that the scheduler no longer deems necessary. For the single-machine case, the method is used less often.

Let's move on to the GPU accelerated version with dask_cudf.

We can use `nvidia-smi` command to check the usage of our GPU.

```
[9]: !nvidia-smi
```

Mon Sep 21 14:10:27 2020

```
+-----+
| NVIDIA-SMI 440.64.00      Driver Version: 440.64.00      CUDA Version: 10.2      |
+-----+-----+-----+
| GPU  Name           Persistence-M| Bus-Id        Disp.A | Volatile Uncorr. ECC |
| Fan  Temp  Perf  Pwr:Usage/Cap|      Memory-Usage | GPU-Util  Compute M. |
+=====+=====+=====+
|   0   Tesla T4              Off | 00000000:00:1E.0 Off |                    0 |
| N/A   31C    P8      9W /  70W |      0MiB / 15109MiB |      0%      Default |
+-----+-----+-----+

+-----+
| Processes:                                     GPU Memory |
|  GPU       PID    Type    Process name                     Usage      |
+=====+
| No running processes found
+-----+
```

```
[10]: import dask as dask, dask_cudf
      from dask.distributed import Client
      from dask_cuda import LocalCUDACluster
```

```
[11]: cluster = LocalCUDACluster(ip='127.0.0.1',
                                rmm_pool_size="16GB")
      client = Client(cluster)
      client
```

```
[11]: <Client: 'tcp://127.0.0.1:34691' processes=1 threads=1, memory=16.52 GB>
```

We reserve 14GB per GPU via rmm_pool_size.

```
[12]: !nvidia-smi
```

Mon Sep 21 14:10:33 2020

```
+-----+
| NVIDIA-SMI 440.64.00      Driver Version: 440.64.00      CUDA Version: 10.2      |
+-----+-----+-----+
| GPU  Name           Persistence-M| Bus-Id        Disp.A | Volatile Uncorr. ECC |
| Fan  Temp  Perf  Pwr:Usage/Cap|      Memory-Usage | GPU-Util  Compute M. |
+=====+=====+=====+
|   0   Tesla T4              Off | 00000000:00:1E.0 Off |                    0 |
| N/A   33C    P0     26W /  70W |     605MiB / 15109MiB |      0%      Default |
+-----+-----+-----+

+-----+
| Processes:                                     GPU Memory |
```

	GPU	PID	Type	Process name	Usage
=====					
+-----+-----+-----+-----+-----+-----+					

We use `dask_cudf` to read and load the data. The remaining code is exactly the same as the `dask pandas` version.

```
[13]: %%time

ddf_train = dask_cudf.read_parquet('./data/train.parquet')
ddf_valid = dask_cudf.read_parquet('./data/valid.parquet')
```

CPU times: user 996 ms, sys: 416 ms, total: 1.41 s
Wall time: 1.41 s

```
[14]: %%time

ddf_train['cat_2_brand'] = ddf_train['cat_2'].astype(str) + '_' +
    ↪ddf_train['brand'].astype(str)
ddf_valid['cat_2_brand'] = ddf_valid['cat_2'].astype(str) + '_' +
    ↪ddf_valid['brand'].astype(str)

ddf_train_group = ddf_train[['cat_2_brand', 'target']].groupby(['cat_2_brand']).
    ↪agg(['count', 'mean'])
ddf_train_group = ddf_train_group.reset_index()
ddf_train_group.columns = ['cat_2_brand', 'TE_count', 'TE_mean']
ddf_train = ddf_train.merge(ddf_train_group, how='left', on='cat_2_brand')
ddf_valid = ddf_valid.merge(ddf_train_group, how='left', on='cat_2_brand')
global_mean = ddf_train['target'].mean()
ddf_train['TE_mean'] = ddf_train.TE_mean.where(ddf_train['TE_count']>20,
    ↪global_mean)
ddf_valid['TE_mean'] = ddf_valid.TE_mean.where(ddf_valid['TE_count']>20,
    ↪global_mean)
```

CPU times: user 544 ms, sys: 0 ns, total: 544 ms
Wall time: 652 ms

```
[15]: %%time

ddf_train.compute()
ddf_valid.compute()
```

CPU times: user 2.75 s, sys: 5.77 s, total: 8.52 s
Wall time: 14.3 s

```
[15]:
```

	event_time	event_type	product_id	brand	price	\
0	2020-03-01 08:16:04 UTC	cart	1005135	apple	1516.10	
1	2020-03-01 08:16:08 UTC	cart	1005135	apple	1516.10	

2	2020-03-01 08:16:09 UTC	cart	1004996	doogee	96.89
3	2020-03-01 08:16:09 UTC	cart	1005135	apple	1516.10
4	2020-03-01 08:16:13 UTC	cart	1005256	xiaomi	141.29
...
2461714	2020-03-31 19:25:14 UTC	purchase	18301044	<NA>	11.04
2461715	2020-03-31 19:25:17 UTC	purchase	100058915	iqos	43.76
2461716	2020-03-31 19:25:36 UTC	purchase	32401283	<NA>	22.97
2461717	2020-03-31 19:26:18 UTC	purchase	4800282	samsung	38.59
2461718	2020-03-31 19:26:20 UTC	purchase	100058915	iqos	43.76

	user_id	user_session	target	\
0	620967403	2f69a6e0-3a9e-4b7c-b717-ce5b8ad85ce3	0	
1	620967403	2f69a6e0-3a9e-4b7c-b717-ce5b8ad85ce3	0	
2	607174356	80d6850c-7f95-4978-ba1a-dedbe802e012	0	
3	620967403	2f69a6e0-3a9e-4b7c-b717-ce5b8ad85ce3	0	
4	571788375	da050faa-118a-405a-b9c8-63f9d730328e	0	
...
2461714	572119027	172b36e9-9259-423c-bc43-5d555ff94ce4	1	
2461715	620477097	47786b4a-f2c3-48fa-b714-9d05556d5b98	1	
2461716	635102002	d82b8bf0-dea5-4e53-84f8-e61332eb17f1	1	
2461717	622434648	4894c1b9-d00d-4418-b6c9-e8cd2f842b33	1	
2461718	620477097	47786b4a-f2c3-48fa-b714-9d05556d5b98	1	

	cat_0	cat_1	...	timestamp	ts_hour	ts_minute	\
0	construction	tools	...	2020-03-01 08:16:04	8	16	
1	construction	tools	...	2020-03-01 08:16:08	8	16	
2	construction	tools	...	2020-03-01 08:16:09	8	16	
3	construction	tools	...	2020-03-01 08:16:09	8	16	
4	construction	tools	...	2020-03-01 08:16:13	8	16	
...
2461714	sport	ski	...	2020-03-31 19:25:14	19	25	
2461715	apparel	trousers	...	2020-03-31 19:25:17	19	25	
2461716	apparel	underwear	...	2020-03-31 19:25:36	19	25	
2461717	sport	bicycle	...	2020-03-31 19:26:18	19	26	
2461718	apparel	trousers	...	2020-03-31 19:26:20	19	26	

	ts_weekday	ts_day	ts_month	ts_year	cat_2_brand	TE_count	\
0	6	1	3	2020	light_apple	1013391	
1	6	1	3	2020	light_apple	1013391	
2	6	1	3	2020	light_doogee	769	
3	6	1	3	2020	light_apple	1013391	
4	6	1	3	2020	light_xiaomi	510657	
...
2461714	1	31	3	2020	<NA>	<NA>	
2461715	1	31	3	2020	<NA>	<NA>	
2461716	1	31	3	2020	<NA>	<NA>	
2461717	1	31	3	2020	<NA>	<NA>	

```
2461718      1      31      3      2020      <NA>      <NA>
```

```
      TE_mean
0      0.469441
1      0.469441
2      0.405722
3      0.469441
4      0.396346
...      ...
2461714 0.366924
2461715 0.366924
2461716 0.366924
2461717 0.366924
2461718 0.366924
```

```
[2461719 rows x 22 columns]
```

```
[16]: !nvidia-smi
```

```
Mon Sep 21 14:10:50 2020
```

```
+-----+
| NVIDIA-SMI 440.64.00      Driver Version: 440.64.00      CUDA Version: 10.2      |
+-----+-----+-----+-----+
| GPU  Name           Persistence-M| Bus-Id        Disp.A | Volatile Uncorr. ECC |
| Fan  Temp   Perf   Pwr:Usage/Cap|      Memory-Usage | GPU-Util  Compute M. |
+-----+-----+-----+-----+
|   0   Tesla T4              Off | 00000000:00:1E.0 Off |                    0 |
| N/A   36C    P0      33W /  70W |  1629MiB / 15109MiB |      0%      Default |
+-----+-----+-----+-----+

+-----+
| Processes:                                     GPU Memory |
|  GPU       PID    Type    Process name                     Usage      |
+-----+-----+-----+-----+
|
```

```
[17]: client.close()
```